

# Gauging Sentiment of Emerging Technology on Twitter



## Business Problem

Gauging sentiment of a market sector, company, or product is very helpful to investment analysis. An investment firm is analyzing emerging technology and would like a predictive model that can analyze text data to classify the sentiment; positive or negative. Being able to gauge overall sentiment will help in analysis of the sector and help them strategically place investments in the market.

## Data Understanding

The dataset consists of thousands of tweets from the SXSW festival pertaining to Apple and Google products. The tweets are labeled as positive, negative, no opinion, or "I can't tell." We will analyze the data from a binary classification standpoint and only keep the positive and negative classes.

The data is heavily imbalanced with positive tweets outnumbering the negative by a 6:1 ratio.

## Methods

Using NLP methods we will pre-process the data to get it ready for modeling.



1. First, we will clean the data by removing stop words, punctuation and other pieces of text that do not add value to the analysis such as numbers, and twitter slang.
2. Then we will tokenize the text with regular expressions
3. Next we will both stem and lemmatize the text separately so that we can train models on both and see which is better.
4. Finally we will vectorize the text with both TF-IDF and Count vectorization so that we again can see which will provide a better model.
5. We will also look at removing mutually exclusive words from our training set to see if it improves our model.

We will split our dataset into training and testing data, and to handle the imbalance we will use SMOTE to provide a balanced training set.

Once the data is ready for modeling we will build a baseline and then train several models with differing parameters, score them with cross validation, and generate predictions from our models to compare with the test set and attain our final results.

```
In [1]: 1 # Import our dependencies
2 import pandas as pd
3 import numpy as np
4 from mpl_toolkits.mplot3d import Axes3D
5 import matplotlib.pyplot as plt
6 %matplotlib inline
7 from sklearn.utils import class_weight
8 from sklearn.model_selection import train_test_split
9 from sklearn.feature_extraction.text import TfidfVectorizer
10 from sklearn.naive_bayes import MultinomialNB
11 from sklearn.model_selection import cross_val_score
12 from sklearn.manifold import TSNE
13 from nltk.tokenize import word_tokenize
14 from nltk.tokenize import RegexpTokenizer
15 from imblearn.over_sampling import SMOTE
16 import string
17 from nltk import FreqDist
18 from nltk.corpus import stopwords
19 from nltk.stem import SnowballStemmer
20 from nltk.stem import WordNetLemmatizer
21 from sklearn.metrics import classification_report
22 from sklearn.metrics import plot_confusion_matrix
23 import matplotlib.pyplot as plt
24 from sklearn.svm import SVC
25 from collections import Counter
26 from imblearn.under_sampling import NearMiss
27 from sklearn.neighbors import KNeighborsClassifier
28 from sklearn.tree import DecisionTreeClassifier
29 from sklearn.ensemble import RandomForestClassifier
30 from sklearn.feature_extraction.text import CountVectorizer
31 import warnings
32 warnings.simplefilter(action='ignore', category=FutureWarning)
```

started 08:23:14 2022-02-21, finished in 1.02s

# EDA

In [2]:

```
1 # Read in our dataset and view the first 5 rows
2 df = pd.read_csv('data/judge_1377884607_tweet_product_company.csv')
3 df.head()
```

started 08:23:15 2022-02-21, finished in 29ms

Out[2]:

	tweet_text	emotion_in_tweet_is_directed_at	is_there_an_emotion_directed_at_a_brand_or_produ
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone	Negative err
1	@jessedee Know about @fludapp ? Awesome iPad/i...	iPad or iPhone App	Positive err
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad	Positive err
3	@sxsxw I hope this year's festival isn't as cra...	iPad or iPhone App	Negative err
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google	Positive err

In [3]:

```
1 # View our columns
2 df.info()
```

started 08:23:15 2022-02-21, finished in 9ms

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8721 entries, 0 to 8720
Data columns (total 3 columns):
#   Column                                Non-Null Count
Dtype
---  -----
0    tweet_text                          8720 non-null
object
1    emotion_in_tweet_is_directed_at     3169 non-null
object
2    is_there_an_emotion_directed_at_a_brand_or_product 8721 non-null
object
dtypes: object(3)
memory usage: 204.5+ KB
```

```
In [4]: 1 # See the various products and companies the tweets are about
        2 df.emotion_in_tweet_is_directed_at.value_counts()
```

started 08:23:15 2022-02-21, finished in 4ms

```
Out[4]: iPad          910
        Apple         640
        iPad or iPhone App  451
        Google        412
        iPhone        288
        Other Google product or service  282
        Android App    78
        Android        74
        Other Apple product or service  34
        Name: emotion_in_tweet_is_directed_at, dtype: int64
```

```
In [5]: 1 # Get a feel for the class labels
        2 df.is_there_an_emotion_directed_at_a_brand_or_product.value_counts()
```

started 08:23:15 2022-02-21, finished in 4ms

```
Out[5]: No emotion toward brand or product    5156
        Positive emotion                    2869
        Negative emotion                    545
        I can't tell                        151
        Name: is_there_an_emotion_directed_at_a_brand_or_product, dtype: int64
```

## Data Cleaning

```
In [6]: 1 # Since we are analyzing sentiment we will drop the labels with no sentiment
        2 pos_neg = df[df['is_there_an_emotion_directed_at_a_brand_or_product'] != 0]
        3 pos_neg = pos_neg[pos_neg['is_there_an_emotion_directed_at_a_brand_or_p
```

started 08:23:15 2022-02-21, finished in 6ms

Extremely unbalanced dataset

```
In [7]: 1 pos_neg.is_there_an_emotion_directed_at_a_brand_or_product.value_counts()
```

started 08:23:15 2022-02-21, finished in 6ms

```
Out[7]: Positive emotion    0.840363
        Negative emotion    0.159637
        Name: is_there_an_emotion_directed_at_a_brand_or_product, dtype: float64
```

```
In [8]: 1 # Drop the products/companies column since we will not be using this in
2 # Change column names to something short and easier to work with
3 pos_neg = pos_neg.drop('emotion_in_tweet_is_directed_at', axis=1)
4 pos_neg = pos_neg.rename(columns={'tweet_text': 'text', 'is_there_an_emo
5 pos_neg.head()
```

started 08:23:15 2022-02-21, finished in 7ms

Out[8]:

	text	target
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	Negative emotion
1	@jessedee Know about @fludapp ? Awesome iPad/i...	Positive emotion
2	@swonderlin Can not wait for #iPad 2 also. The...	Positive emotion
3	@sxsxw I hope this year's festival isn't as cra...	Negative emotion
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Positive emotion

```
In [9]: 1 # Create a function to remove Twitter lingo like @'s and #'s and map it
2 def remove_ats_and_hashtags(text):
3     entity_prefixes = ['@', '#', '0']
4     for separator in string.punctuation:
5         if separator not in entity_prefixes :
6             text = text.replace(separator, ' ')
7     words = []
8     for word in text.split():
9         word = word.strip()
10        if word:
11            if word[0] not in entity_prefixes:
12                words.append(word)
13    return ' '.join(words)
14
15 pos_neg['text'] = pos_neg['text'].map(remove_ats_and_hashtags)
```

started 08:23:15 2022-02-21, finished in 38ms

```
In [10]: 1 # Encode our classes/targets for modeling
2 pos_neg.replace({'Negative emotion' : 0, 'Positive emotion' : 1}, inplace=True)
```

started 08:23:15 2022-02-21, finished in 4ms

```
In [11]: 1 # Transform our text to all lower case letters
2 pos_neg['text'] = pos_neg['text'].str.lower()
3 pos_neg.head()
```

started 08:23:15 2022-02-21, finished in 7ms

Out[11]:

	text	target
0	i have a 3g iphone after 3 hrs tweeting at aus...	0
1	know about awesome ipad iphone app that you ll...	1
2	can not wait for 2 also they should sale them ...	1
3	i hope this year s festival isn t as crashy as...	0
4	great stuff on fri marissa mayer google tim o ...	1

## Explore the words

```
In [12]: 1 # Tokenize our text to split into words
2 pos_neg['text_tokenized'] = pos_neg['text'].apply(word_tokenize)
```


started 08:23:15 2022-02-21, finished in 334ms

```
In [13]: 1 # Create a function that will view the 20 most frequently used words
2 def visualize_top_20(freq_dist, title):
3
4     # Extract data for plotting
5     top_20 = list(zip(*freq_dist.most_common(20)))
6     tokens = top_20[0]
7     counts = top_20[1]
8
9     # Set up plot and plot data
10    fig, ax = plt.subplots(figsize=(20,10))
11    ax.bar(tokens, counts)
12
13    # Customize plot appearance
14    ax.set_title(title)
15    ax.set_ylabel("Count")
16    ax.tick_params(axis="x", rotation=90)
```

started 08:23:15 2022-02-21, finished in 4ms

```
In [14]: 1 # create a frequency distribution and view it as a histogram
2 pos_neg_freq_dist = FreqDist(pos_neg['text_tokenized'].explode())
3
4 visualize_top_20(pos_neg_freq_dist, "Top 20 Word Frequency")
```

started 08:23:15 2022-02-21, finished in 220ms



There are a number of stopwords, 1 character and other phrases that won't add value to our model.

## Remove stopwords, punctuation and other unwanted phrases

```
In [15]: 1 # Create a stopwords list for removal, add punctuation and other phrase
2 stopwords_list = stopwords.words('english')
3 stopwords_list += list(string.punctuation)
4 new_stops = ('quot', 'rt', 'i', 'amp')
5 stopwords_list += list(new_stops)
```

started 08:23:15 2022-02-21, finished in 4ms

```
In [16]: 1 # Create a function that will remove the stopwords list from the text a
2 def remove_stopwords(token_list):
3     """
4     Given a list of tokens, return a list where the tokens
5     that are in stopwords_list have been removed
6     """
7     stops_rmv_list = [token for token in token_list if token not in sto
8     return stops_rmv_list
9
10 pos_neg['stopwords_removed'] = pos_neg['text_tokenized'].apply(remove_s
```

started 08:23:15 2022-02-21, finished in 100ms

```
In [17]: 1 # Remove numbers and single letters and create a text column for modeli
2 pattern = "([a-z]{4,})"
3 regex_tokenizer = RegexpTokenizer(pattern)
4
5 pos_neg['regex_text'] = [' '.join(text) for text in pos_neg.stopwords_r
6 pos_neg['regex_text_tokenized'] = [regex_tokenizer.tokenize(text) for t
7 pos_neg['regex_text'] = [' '.join(text) for text in pos_neg.regex_text_
started 08:23:15 2022-02-21, finished in 19ms
```

```
In [18]: 1 pos_neg.head( )
started 08:23:15 2022-02-21, finished in 11ms
```

Out[18]:

	text	target	text_tokenized	stopwords_removed	regex_text	regex_text_tokenized
0	i have a 3g iphone after 3 hrs tweeting at aus...	0	[i, have, a, 3g, iphone, after, 3, hrs, tweeti...	[3g, iphone, 3, hrs, tweeting, austin, dead, n...	iphone tweeting austin dead need upgrade plugi...	[iphone, tweeting, austin, dead, need, upgrade...
1	know about awesome ipad iphone app that you ll...	1	[know, about, awesome, ipad, iphone, app, that...	[know, awesome, ipad, iphone, app, likely, app...	know awesome ipad iphone likely appreciate des...	[know, awesome, ipad, iphone, likely, apprecia...
2	can not wait for 2 also they should sale them ...	1	[can, not, wait, for, 2, also, they, should, s...	[wait, 2, also, sale]	wait also sale	[wait, also, sale]
3	i hope this year s festival isn t as crashy as...	0	[i, hope, this, year, s, festival, isn, t, as,...	[hope, year, festival, crashy, year, iphone, app]	hope year festival crashy year iphone	[hope, year, festival, crashy, year, iphone]
4	great stuff on fri marissa mayer google tim o ...	1	[great, stuff, on, fri, marissa, mayer, google...	[great, stuff, fri, marissa, mayer, google, ti...	great stuff marissa mayer google reilly tech b...	[great, stuff, marissa, mayer, google, reilly,...

Stem words



```
In [19]: 1 snow_stemmer = SnowballStemmer(language="english")
2
3 pos_neg['stemmed_text'] = [snow_stemmer.stem(text) for text in pos_neg[
4
5 pos_neg.head()]

started 08:23:15 2022-02-21, finished in 70ms
```

Out[19]:

	text	target	text_tokenized	stopwords_removed	regex_text	regex_text_tokenized
0	i have a 3g iphone after 3 hrs tweeting at aus...	0	[i, have, a, 3g, iphone, after, 3, hrs, tweeti...	[3g, iphone, 3, hrs, tweeting, austin, dead, n...	iphone tweeting austin dead need upgrade plugi...	[iphone, tweeting, austin, dead, need, upgrade...
1	know about awesome ipad iphone app that you ll...	1	[know, about, awesome, ipad, iphone, app, that...	[know, awesome, ipad, iphone, app, likely, app...	know awesome ipad iphone likely appreciate des...	[know, awesome, ipad, iphone, likely, apprecia...
2	can not wait for 2 also they should sale them ...	1	[can, not, wait, for, 2, also, they, should, s...	[wait, 2, also, sale]	wait also sale	[wait, also, sale]
3	i hope this year s festival isn t as crashy as...	0	[i, hope, this, year, s, festival, isn, t, as,...	[hope, year, festival, crashy, year, iphone, app]	hope year festival crashy year iphone	[hope, year, festival, crashy, year, iphone]
4	great stuff on fri marissa mayer google tim o ...	1	[great, stuff, on, fri, marissa, mayer, google...	[great, stuff, fri, marissa, mayer, google, ti...	great stuff marissa mayer google reilly tech b...	[great, stuff, marissa, mayer, google, reilly,...

## Lemmatize words

```
In [20]: 1 lemmer = WordNetLemmatizer()
2
3 pos_neg['lemmed_text'] = [lemmer.lemmatize(text) for text in pos_neg['r
4
5 pos_neg.head(10)
```

started 08:23:15 2022-02-21, finished in 1.28s

Out[20]:

	text	target	text_tokenized	stopwords_removed	regex_text	regex_text_tokenized
0	i have a 3g iphone after 3 hrs tweeting at aus...	0	[i, have, a, 3g, iphone, after, 3, hrs, tweeti...	[3g, iphone, 3, hrs, tweeting, austin, dead, n...	iphone tweeting austin dead need upgrade plugi...	[iphone, tweeting, austin, dead, need, upgrade.
1	know about awesome ipad iphone app that you ll...	1	[know, about, awesome, ipad, iphone, app, that...	[know, awesome, ipad, iphone, app, likely, app...	know awesome ipad iphone likely appreciate des...	[know, awesome, ipad, iphone, likely, apprecia.
2	can not wait for 2 also they should sale them ...	1	[can, not, wait, for, 2, also, they, should, s...	[wait, 2, also, sale]	wait also sale	[wait, also, sale]
3	i hope this year s festival isn t as crashy as...	0	[i, hope, this, year, s, festival, isn, t, as,...	[hope, year, festival, crashy, year, iphone, app]	hope year festival crashy year iphone	[hope, year, festival, crashy, year, iphone]
4	great stuff on fri marissa mayer google tim o ...	1	[great, stuff, on, fri, marissa, mayer, google...	[great, stuff, fri, marissa, mayer, google, ti...	great stuff marissa mayer google reilly tech b...	[great, stuff, marissa, mayer, google, reilly..
7	is just starting is around the corner and is o...	1	[is, just, starting, is, around, the, corner, ...	[starting, around, corner, hop, skip, jump, go...	starting around corner skip jump good time	[starting, around, corner, skip, jump, good, t.
8	beautifully smart and simple idea rt wrote abo...	1	[beautifully, smart, and, simple, idea, rt, wr...	[beautifully, smart, simple, idea, wrote, ipad...	beautifully smart simple idea wrote ipad http ...	[beautifully, smart, simple, idea, wrote, ipad.

	text	target	text_tokenized	stopwords_removed	regex_text	regex_text_tokenized
9	counting down the days to plus strong canadian...	1	[counting, down, the, days, to, plus, strong, ...	[counting, days, plus, strong, canadian, dolla...	counting days plus strong canadian dollar mean...	[counting, days, plu: strong, canadian, dolla.
10	excited to meet the at so i can show them my s...	1	[excited, to, meet, the, at, so, i, can, show,...	[excited, meet, show, sprint, galaxy, still, r...	excited meet show sprint galaxy still running ...	[excited, meet, shov sprint, galaxy, still, r.
11	find amp start impromptu parties at with http ...	1	[find, amp, start, impromptu, parties, at, wit...	[find, start, impromptu, parties, http, bit, l...	find start impromptu parties http gvlrin wait ...	[find, start, imprompti parties, http, gvlrin.


You can see in several of the tweets that stemming removes the 'e' from words like 'sale' and 'iphone' so we will use the lemmatized words for modeling

## View the new frequency distribution of words after processing

```
In [21]: 1 pos_neg['lemmed_text_tokenized'] = pos_neg['lemmed_text'].apply(word_to
started 08:23:17 2022-02-21, finished in 273ms
```

```
In [22]: 1 # create a frequency distribution and view it as a histogram
2 lemmed_text_freq_dist = FreqDist(pos_neg['lemmed_text_tokenized']).explo
3
4 visualize_top_20(lemmed_text_freq_dist, "Top 20 Word Frequency")
```

started 08:23:17 2022-02-21, finished in 190ms



## Train Test Split

```
In [23]: 1 X = pos_neg.drop(['target'], axis=1)
2 y = pos_neg['target']
3
4 X_train , X_test, y_train, y_test = train_test_split(
5     X, y, test_size=0.2, random_state=30)
```

started 08:23:17 2022-02-21, finished in 8ms

```
In [24]: 1 y_train.value_counts(normalize=True)
```

started 08:23:17 2022-02-21, finished in 5ms

```
Out[24]: 1    0.844746
0    0.155254
Name: target, dtype: float64
```

## Build and score a baseline model

### Using Count Vectorization on lemmatized text

```
In [25]: 1 # Use count vectorization
2 count = CountVectorizer()
3 X_count_vectorized = count.fit_transform(X_train['lemmed_text'])
```

started 08:23:17 2022-02-21, finished in 28ms

```
In [26]: 1 # With such an imbalanced dataset use SMOTE to balance training data
2 smote_count = SMOTE(k_neighbors=3)
3 X_train_resampled, y_train_resampled = smote_count.fit_resample(
4     X_count_vectorized, y_train)
```

started 08:23:17 2022-02-21, finished in 14ms

```
In [27]: 1 # Check that the training set is balanced now
2 y_train_resampled.value_counts(normalize=True)
```

started 08:23:17 2022-02-21, finished in 5ms

```
Out[27]: 1    0.5
0    0.5
Name: target, dtype: float64
```

```
In [28]: 1 # Instantiate a MultinomialNB classifier and fi it to training data
2 baseline_model = MultinomialNB()
3
4 baseline_model.fit(X_train_resampled, y_train_resampled)
5
6 # Evaluate the model
7 baseline_cv = cross_val_score(baseline_model, X_train_resampled, y_train
8 print("Baseline:", baseline_cv.mean()))
```

started 08:23:17 2022-02-21, finished in 19ms

Baseline: 0.7986608790067286

```
In [29]: 1 # Use count vectorization with bigrams
2 count_bigram = CountVectorizer(ngram_range=(1,2))
3 X_count_vectorized_bigrams = count_bigram.fit_transform(X_train['lemmed
```

started 08:23:17 2022-02-21, finished in 61ms

```
In [30]: 1 smote_count = SMOTE(k_neighbors=3)
2 X_train_resampled_bigram, y_train_resampled_bigram = smote_count.fit_re
3     X_count_vectorized_bigrams, y_train)
```

started 08:23:17 2022-02-21, finished in 15ms

```
In [31]: 1 baseline_model_bigram = MultinomialNB()
2
3 baseline_model_bigram.fit(X_train_resampled_bigram, y_train_resampled_b
4
5 baseline_bigram_cv = cross_val_score(baseline_model_bigram, X_train_res
6 print("Baseline          :", baseline_cv.mean())
7 print("Baseline w/ bigrams :", baseline_bigram_cv.mean())
```

started 08:23:17 2022-02-21, finished in 21ms

Baseline : 0.7986608790067286  
Baseline w/ bigrams : 0.7373374570802085

## Score using TF-IDF Vectorization

```
In [32]: 1 tfidf = TfidfVectorizer(max_features=20)
2 X_train_vectorized_tfidf = tfidf.fit_transform(X_train['lemmed_text'])
```

started 08:23:17 2022-02-21, finished in 31ms

```
In [33]: 1 smote_count = SMOTE(k_neighbors=3)
2 X_train_resampled_tfidf, y_train_resampled_tfidf = smote_count.fit_resa
3 X_train_vectorized_tfidf, y_train)
```

started 08:23:17 2022-02-21, finished in 9ms

```
In [34]: 1 baseline_model_tfidf = MultinomialNB()
2
3 baseline_model_tfidf.fit(X_train_resampled_tfidf, y_train_resampled_tfi
4
5 baseline_tfidf_cv = cross_val_score(baseline_model_tfidf, X_train_resam
6 print("Baseline          :", baseline_cv.mean())
7 print("Baseline w/ bigrams :", baseline_bigram_cv.mean())
8 print("Baseline TF-IDF    :", baseline_tfidf_cv.mean())
```

started 08:23:17 2022-02-21, finished in 18ms

```
Baseline          : 0.7986608790067286
Baseline w/ bigrams : 0.7373374570802085
Baseline TF-IDF    : 0.6569201627250572
```


**Looks like the regular baseline model scored the best - "Multinomial Naive Bayes"**

Lets see how it does on test data

```
In [35]: 1 X_test_vectorized = count.transform(X_test['lemmed_text'])
2 baseline_preds = baseline_model.predict(X_test_vectorized)
3
4 print(classification_report(y_test, baseline_preds))
5
6 plot_confusion_matrix(baseline_model, X_test_vectorized, y_test, cmap=p
7 plt.grid(False)
8 plt.show()
```

started 08:23:17 2022-02-21, finished in 130ms

	precision	recall	f1-score	support
0	0.54	0.60	0.56	121
1	0.91	0.89	0.90	562
accuracy			0.84	683
macro avg	0.72	0.74	0.73	683
weighted avg	0.84	0.84	0.84	683



**This scores well**

Overall Accuracy of 84%


Predicts positive sentiment with an accuracy of 89%  
Predicts negative sentiment with an accuracy of 61%

**And the bigrams model too**

```
In [36]: 1 X_test_vectorized_bigrams = count_bigram.transform(X_test['lemmed_text'])
2 baseline_preds = baseline_model_bigram.predict(X_test_vectorized_bigram)
3
4 bigram_preds = baseline_model_bigram.predict(X_test_vectorized_bigrams)
5
6 print(classification_report(y_test, bigram_preds))
7
8 plot_confusion_matrix(baseline_model_bigram, X_test_vectorized_bigrams,
9 plt.grid(False)
10 plt.show())
```

started 08:23:18 2022-02-21, finished in 136ms

	precision	recall	f1-score	support
0	0.64	0.56	0.60	121
1	0.91	0.93	0.92	562
accuracy			0.87	683
macro avg	0.77	0.75	0.76	683
weighted avg	0.86	0.87	0.86	683



### Bigrams also score well

Overall Accuracy of 86%

Predicts positive sentiment with an accuracy of 93%  
Predicts negative sentiment with an accuracy of 55%

**Lets run the model on the original text with no processing and see how it does**

```
In [37]: 1 count_og_text = CountVectorizer()
2 X_count_og_text_vectorized = count_og_text.fit_transform(X_train['text'])
```

started 08:23:18 2022-02-21, finished in 40ms



```
In [38]: 1 smote_og_count = SMOTE(k_neighbors=3)
2 X_train_resampled_og, y_train_resampled_og = smote_og_count.fit_resample(
3     X_count_og_text_vectorized, y_train)
```

started 08:23:18 2022-02-21, finished in 16ms

```
In [39]: 1 baseline_model_og = MultinomialNB()
2
3 baseline_model_og.fit(X_train_resampled_og, y_train_resampled_og)
4
5 # Evaluate the model
6 baseline_og_cv = cross_val_score(baseline_model_og, X_train_resampled_og,
7     y_train_resampled_og, cv=5)
8 print("Baseline : ", baseline_cv.mean())
9 print("Baseline w/ og text : ", baseline_og_cv.mean())
```


started 08:23:18 2022-02-21, finished in 22ms

```
Baseline : 0.7986608790067286
Baseline w/ og text : 0.8129674761400038
```

```
In [40]: 1 X_test_og_vectorized = count_og_text.transform(X_test['text'])
2 baseline_og_preds = baseline_model_og.predict(X_test_og_vectorized)
3
4 print(classification_report(y_test, baseline_og_preds))
5
6 plot_confusion_matrix(baseline_model_og, X_test_og_vectorized, y_test,
7     plt.grid(False))
8 plt.show()
```

started 08:23:18 2022-02-21, finished in 115ms

	precision	recall	f1-score	support
0	0.53	0.49	0.51	121
1	0.89	0.91	0.90	562
accuracy			0.83	683
macro avg	0.71	0.70	0.70	683
weighted avg	0.83	0.83	0.83	683



Surprisingly it scores almost as well as the model with lemmatized text

Overall Accuracy of 83%

Predicts positive sentiment with an accuracy of 91%  
Predicts negative sentiment with an accuracy of 50%

**We will stick with the Lemmatized text model**

**Let's run a few more models to see if we can score better**

We will stick with these parameters:

Lemmatized text  
Count vectorization

```
In [41]: 1 # K Nearest Neighbors with n=5
2 knn_5 = KNeighborsClassifier(n_neighbors=5)
3 knn_5.fit(X_train_resampled, y_train_resampled)
4
5 knn_5_cv = cross_val_score(knn_5, X_train_resampled, y_train_resampled)
6 print("Baseline :", baseline_cv.mean())
7 print("KNN_5    :", knn_5_cv.mean())
```


started 08:23:18 2022-02-21, finished in 392ms

Baseline : 0.7986608790067286  
KNN\_5 : 0.5806259885359211

```
In [42]: 1 knn_5_preds = knn_5.predict(X_test_vectorized)
2
3 print(classification_report(y_test, knn_5_preds))
4
5 plot_confusion_matrix(knn_5, X_test_vectorized, y_test, cmap=plt.cm.Ora
6 plt.grid(False)
7 plt.show()
```

started 08:23:18 2022-02-21, finished in 245ms

	precision	recall	f1-score	support
0	0.21	0.98	0.34	121
1	0.97	0.20	0.33	562
accuracy			0.33	683
macro avg	0.59	0.59	0.33	683
weighted avg	0.84	0.33	0.33	683



```
In [43]: 1 # K Nearest Neighbors with n=3
2 knn_3 = KNeighborsClassifier(n_neighbors=3)
3 knn_3.fit(X_train_resampled, y_train_resampled)
4
5 knn_3_cv = cross_val_score(knn_3, X_train_resampled, y_train_resampled)
6 print("Baseline :", baseline_cv.mean())
7 print("KNN_5    :", knn_5_cv.mean())
8 print("KNN_3    :", knn_3_cv.mean())
```


started 08:23:19 2022-02-21, finished in 312ms

```
Baseline : 0.7986608790067286
KNN_5    : 0.5806259885359211
KNN_3    : 0.6896421411834934
```

```
In [44]: 1 knn_3_preds = knn_3.predict(X_test_vectorized)
2
3 print(classification_report(y_test, knn_3_preds))
4
5 plot_confusion_matrix(knn_3, X_test_vectorized, y_test, cmap=plt.cm.Oran
6 plt.grid(False)
7 plt.show()
```

started 08:23:19 2022-02-21, finished in 208ms

	precision	recall	f1-score	support
0	0.25	0.88	0.38	121
1	0.94	0.42	0.58	562
accuracy			0.50	683
macro avg	0.59	0.65	0.48	683
weighted avg	0.82	0.50	0.55	683



```
In [45]: 1 # Decision Tree (default)
2 tree = DecisionTreeClassifier()
3 tree.fit(X_train_resampled, y_train_resampled)
4
5 tree_cv = cross_val_score(tree, X_train_resampled, y_train_resampled)
6 print("Baseline :", baseline_cv.mean())
7 print("KNN_5    :", knn_5_cv.mean())
8 print("KNN_3    :", knn_3_cv.mean())
9 print("Tree      :", tree_cv.mean())
```


started 08:23:19 2022-02-21, finished in 453ms

```
Baseline : 0.7986608790067286
KNN_5    : 0.5806259885359211
KNN_3    : 0.6896421411834934
Tree     : 0.8446086161554677
```

```
In [46]: 1 tree_preds = tree.predict(X_test_vectorized)
2
3 print(classification_report(y_test, tree_preds))
4
5 plot_confusion_matrix(tree, X_test_vectorized, y_test, cmap=plt.cm.Oran
6 plt.grid(False)
7 plt.show()
```

started 08:23:20 2022-02-21, finished in 115ms

	precision	recall	f1-score	support
0	0.38	0.60	0.46	121
1	0.90	0.79	0.84	562
accuracy			0.76	683
macro avg	0.64	0.69	0.65	683
weighted avg	0.81	0.76	0.77	683



```
In [47]: 1 # Random Forest (default)
2 forest = RandomForestClassifier()
3 forest.fit(X_train_resampled, y_train_resampled)
4
5 forest_cv = cross_val_score(forest, X_train_resampled, y_train_resample
6 print("Baseline :", baseline_cv.mean())
7 print("KNN_5    :", knn_5_cv.mean())
8 print("KNN_3    :", knn_3_cv.mean())
9 print("Tree     :", tree_cv.mean())
10 print("Forest   :", forest_cv.mean())
```


started 08:23:20 2022-02-21, finished in 10.2s

```
Baseline : 0.7986608790067286
KNN_5    : 0.5806259885359211
KNN_3    : 0.6896421411834934
Tree     : 0.8446086161554677
Forest   : 0.8591309579485926
```

```
In [48]: 1 forest_preds = forest.predict(X_test_vectorized)
2
3 print(classification_report(y_test, forest_preds))
4
5 plot_confusion_matrix(forest, X_test_vectorized, y_test, cmap=plt.cm.Or
6 plt.grid(False)
7 plt.show()
```

started 08:23:30 2022-02-21, finished in 167ms

	precision	recall	f1-score	support
0	0.41	0.56	0.47	121
1	0.90	0.83	0.86	562
accuracy			0.78	683
macro avg	0.65	0.69	0.67	683
weighted avg	0.81	0.78	0.79	683



```
In [49]: 1 # Random Forest w/ bootstrap=False
2 forest_boot = RandomForestClassifier(bootstrap=False)
3 forest_boot.fit(X_train_resampled, y_train_resampled)
4
5 forest_boot_cv = cross_val_score(forest_boot, X_train_resampled, y_train_resampled, cv=5)
6 print("Baseline      :", baseline_cv.mean())
7 print("KNN_5        :", knn_5_cv.mean())
8 print("KNN_3        :", knn_3_cv.mean())
9 print("Tree          :", tree_cv.mean())
10 print("Forest         :", forest_cv.mean())
11 print("Forest Boot    :", forest_boot_cv.mean())
```


started 08:23:30 2022-02-21, finished in 15.8s

```
Baseline      : 0.7986608790067286
KNN_5         : 0.5806259885359211
KNN_3         : 0.6896421411834934
Tree          : 0.8446086161554677
Forest        : 0.8591309579485926
Forest Boot   : 0.8766850057461404
```

```
In [50]: 1 forest_boot_preds = forest_boot.predict(X_test_vectorized)
2
3 print(classification_report(y_test, forest_boot_preds))
4
5 plot_confusion_matrix(forest_boot, X_test_vectorized, y_test, cmap=plt.
6 plt.grid(False)
7 plt.show())
```

started 08:23:46 2022-02-21, finished in 174ms

	precision	recall	f1-score	support
0	0.44	0.50	0.47	121
1	0.89	0.86	0.87	562
accuracy			0.80	683
macro avg	0.66	0.68	0.67	683
weighted avg	0.81	0.80	0.80	683



```
In [51]: 1 # Support Vector Machine (default)
2 clf = SVC()
3 clf.fit(X_train_resampled, y_train_resampled)
4
5 clf_cv = cross_val_score(clf, X_train_resampled, y_train_resampled)
6 print("Baseline      :", baseline_cv.mean())
7 print("KNN_5        :", knn_5_cv.mean())
8 print("KNN_3        :", knn_3_cv.mean())
9 print("Tree         :", tree_cv.mean())
10 print("Forest        :", forest_cv.mean())
11 print("Forest Boot   :", forest_boot_cv.mean())
12 print("SVM          :", clf_cv.mean())
```


started 08:23:46 2022-02-21, finished in 3.54s

Baseline	: 0.7986608790067286
KNN_5	: 0.5806259885359211
KNN_3	: 0.6896421411834934
Tree	: 0.8446086161554677
Forest	: 0.8591309579485926
Forest Boot	: 0.8766850057461404
SVM	: 0.9001012918827834

```
In [52]: 1 clf_preds = clf.predict(X_test_vectorized)
2
3 print(classification_report(y_test, clf_preds))
4
5 plot_confusion_matrix(clf, X_test_vectorized, y_test, cmap=plt.cm.Orang
6 plt.grid(False)
7 plt.show()
```

started 08:23:50 2022-02-21, finished in 292ms

	precision	recall	f1-score	support
0	0.42	0.28	0.34	121
1	0.86	0.92	0.88	562
accuracy			0.80	683
macro avg	0.64	0.60	0.61	683
weighted avg	0.78	0.80	0.79	683




**Out of curiosity, let's see how the Decision Tree, The Random Forrests, and the SVM score with the original text**



```
In [53]: 1 tree_og = DecisionTreeClassifier()
2 tree_og.fit(X_train_resampled_og, y_train_resampled_og)
3
4 tree_og_preds = tree_og.predict(X_test_og_vectorized)
5
6 print(classification_report(y_test, tree_og_preds))
7
8 plot_confusion_matrix(tree_og, X_test_og_vectorized, y_test, cmap=plt.c
9 plt.grid(False)
10 plt.show()
```

started 08:23:50 2022-02-21, finished in 220ms


	precision	recall	f1-score	support
0	0.36	0.45	0.40	121
1	0.87	0.83	0.85	562
accuracy			0.76	683
macro avg	0.62	0.64	0.63	683
weighted avg	0.78	0.76	0.77	683



```
In [54]: 1 forest_og = RandomForestClassifier()
2 forest_og.fit(X_train_resampled_og, y_train_resampled_og)
3
4 forest_og_preds = forest_og.predict(X_test_og_vectorized)
5
6 print(classification_report(y_test, forest_og_preds))
7
8 plot_confusion_matrix(forest_og, X_test_og_vectorized, y_test, cmap=plt
9 plt.grid(False)
10 plt.show())
```

started 08:23:50 2022-02-21, finished in 1.47s


	precision	recall	f1-score	support
0	0.38	0.40	0.39	121
1	0.87	0.86	0.87	562
accuracy			0.78	683
macro avg	0.63	0.63	0.63	683
weighted avg	0.78	0.78	0.78	683



```
In [55]: 1 forest_boot_og = RandomForestClassifier(bootstrap=False)
2 forest_boot_og.fit(X_train_resampled_og, y_train_resampled_og)
3
4 forest_boot_og_preds = forest_boot_og.predict(X_test_og_vectorized)
5
6 print(classification_report(y_test, forest_boot_og_preds))
7
8 plot_confusion_matrix(forest_boot_og, X_test_og_vectorized, y_test, cma
9 plt.grid(False)
10 plt.show())
```

started 08:23:52 2022-02-21, finished in 1.82s


	precision	recall	f1-score	support
0	0.43	0.40	0.41	121
1	0.87	0.89	0.88	562
accuracy			0.80	683
macro avg	0.65	0.64	0.65	683
weighted avg	0.79	0.80	0.80	683



```
In [56]: 1 clf_og = SVC()
2 clf_og.fit(X_train_resampled_og, y_train_resampled_og)
3
4 clf_og_preds = clf_og.predict(X_test_og_vectorized)
5
6 print(classification_report(y_test, clf_og_preds))
7
8 plot_confusion_matrix(clf_og, X_test_og_vectorized, y_test, cmap=plt.cm
9 plt.grid(False)
10 plt.show())
```

started 08:23:53 2022-02-21, finished in 1.82s

	precision	recall	f1-score	support
0	0.42	0.24	0.31	121
1	0.85	0.93	0.89	562
accuracy			0.81	683
macro avg	0.64	0.58	0.60	683
weighted avg	0.77	0.81	0.78	683



**None of these models are as accurate as our Multinomial Naive Bayes models (baseline or bigrams) with lemmatized text**

**Let's do an experiment with mutually exclusive text**

We will remove mutually exclusive text from our training data and see how it affects our models.

**Let's do some more EDA and processing for this new data**

```
In [57]: 1 # Rejoin our training data sets
        2 pos_neg_train = X_train.join(y_train)
```

started 08:23:55 2022-02-21, finished in 4ms

```
In [58]: 1 pos_neg_train.target.value_counts()
```

started 08:23:55 2022-02-21, finished in 6ms

```
Out[58]: 1    2307
        0     424
        Name: target, dtype: int64
```

```
In [59]: 1 # Split the data set into positive and negative for analysis
        2 positive = pos_neg_train.loc[pos_neg['target'] == 1]
        3 negative = pos_neg_train.loc[pos_neg['target'] == 0]
```

started 08:23:55 2022-02-21, finished in 6ms

```
In [60]: 1 positive.info()
```

started 08:23:55 2022-02-21, finished in 12ms

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2307 entries, 1435 to 4378
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   text                  2307 non-null   object
 1   text_tokenized        2307 non-null   object
 2   stopwords_removed     2307 non-null   object
 3   regex_text            2307 non-null   object
 4   regex_text_tokenized  2307 non-null   object
 5   stemmed_text          2307 non-null   object
 6   lemmed_text           2307 non-null   object
 7   lemmed_text_tokenized 2307 non-null   object
 8   target                2307 non-null   int64
dtypes: int64(1), object(8)
memory usage: 180.2+ KB
```

```
In [61]: 1 negative.info()
```

started 08:23:55 2022-02-21, finished in 8ms

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 424 entries, 2997 to 2252
```

```
Data columns (total 9 columns):
```

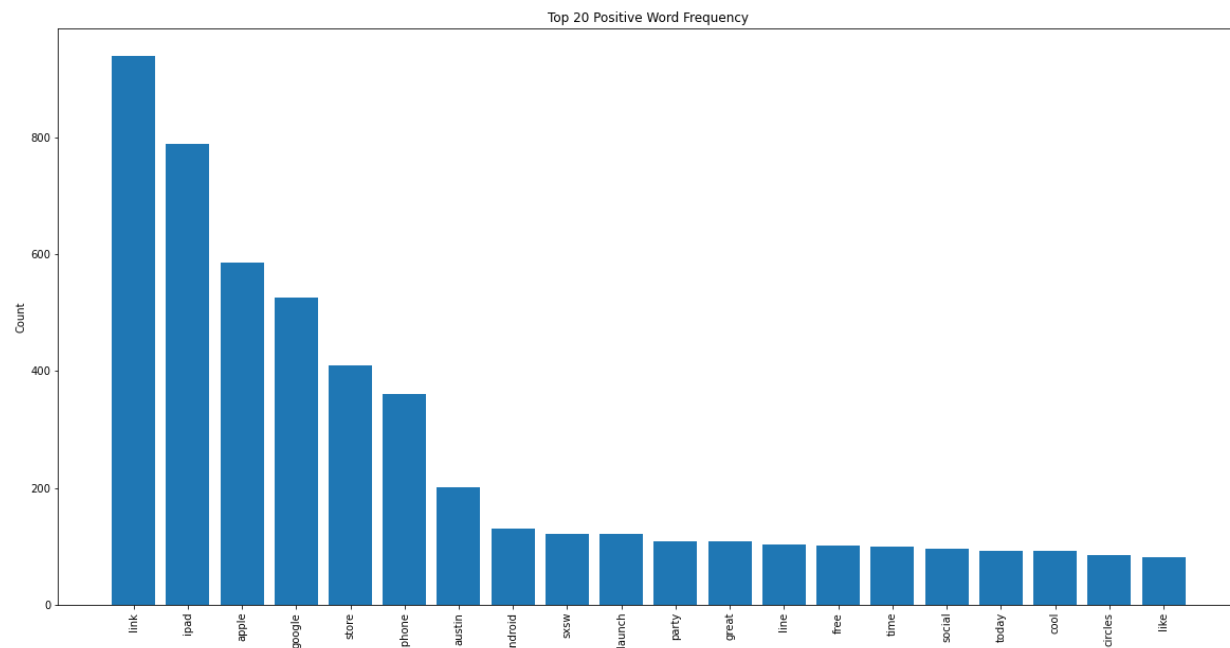
#	Column	Non-Null Count	Dtype
0	text	424 non-null	object
1	text_tokenized	424 non-null	object
2	stopwords_removed	424 non-null	object
3	regex_text	424 non-null	object
4	regex_text_tokenized	424 non-null	object
5	stemmed_text	424 non-null	object
6	lemmed_text	424 non-null	object
7	lemmed_text_tokenized	424 non-null	object
8	target	424 non-null	int64

```
dtypes: int64(1), object(8)
```

```
memory usage: 33.1+ KB
```


```
In [62]: 1 # View the 20 most frequently used words in the positive class
2 positive_freq_dist = FreqDist(positive['lemmed_text_tokenized'].explode()
3
4 visualize_top_20(positive_freq_dist, "Top 20 Positive Word Frequency")
```

started 08:23:55 2022-02-21, finished in 197ms



```
In [63]: 1 # View the 20 most frequently used words in the negative class
2 negative_freq_dist = FreqDist(negative['lemmed_text_tokenized'].explode())
3
4 visualize_top_20(negative_freq_dist, "Top 20 Negative Word Frequency")
```

started 08:23:55 2022-02-21, finished in 169ms



**This is not very telling, and many words are the same, so lets see how removing mutual words affects this**

```
In [64]: 1 # Create a 'Bag of Words' for the positive class
2 BoW_pos = [word for sentence in positive['lemmed_text_tokenized'] for w
```

started 08:23:56 2022-02-21, finished in 4ms

In [65]: 1 BoW\_pos

started 08:23:56 2022-02-21, finished in 13ms

Out[65]: ['wise',  
'apple',  
'opening',  
'temp',  
'store',  
'austin',  
'link',  
'apparently',  
'tell',  
'bizzy',  
'android',  
'remedied',  
'marissa',  
'mayer',  
'connect',  
'digital',  
'physical',  
'worlds',  
'mobile',  
'...']

In [66]: 1 *# Create a 'Bag of Words' for the negative class*  
2 BoW\_neg = [word for sentence in negative['lemmed\_text\_tokenized'] for w

started 08:23:56 2022-02-21, finished in 3ms

In [67]: 1 BoW\_neg

started 08:23:56 2022-02-21, finished in 12ms

Out[67]: ['turning',  
'twitter',  
'forgotten',  
'reason',  
'google',  
'social',  
'technical',  
'dense',  
'vuelta',  
'para',  
'gran',  
'diferencia',  
'revolution',  
'clumsily',  
'translated',  
'google',  
'seems',  
'like',  
'news',  
'...']

In [68]: 1 *# Remove mutual words from the positive BoW*  
2 diff\_neg = [word for word in BoW\_neg if word not in BoW\_pos]

started 08:23:56 2022-02-21, finished in 239ms



```
In [69]: 1 # Remove mutual words from the positive BoW
2 diff_pos = [word for word in BoW_pos if word not in BoW_neg]
```

started 08:23:56 2022-02-21, finished in 346ms

```
In [70]: 1 # Create a 'Bag of Words' that are mutual (appear in both classes)
2 sames = [word for word in BoW_pos if word in BoW_neg]
```

started 08:23:56 2022-02-21, finished in 345ms

```
In [71]: 1 len(sames)
```

started 08:23:57 2022-02-21, finished in 3ms

Out[71]: 12434

```
In [72]: 1 # Create a function that will generate a dictionary of words and their f
2 def counts (lst, series):
3     count_dict = {}
4     for word in lst:
5         count = 0
6         for line in series:
7             if word in line:
8                 count += 1
9                 count_dict[word] = count
10    return count_dict
```

started 08:23:57 2022-02-21, finished in 4ms

```
In [73]: 1 counts(diff_pos, positive['lemmed_text'])
```

started 08:23:57 2022-02-21, finished in 1.47s

```
Out[73]: {'wise': 6,
'bizzy': 4,
'remedied': 1,
'physical': 16,
'worlds': 12,
'soundtrckr': 1,
'featured': 8,
'kick': 9,
'giving': 23,
'visit': 7,
'enter': 28,
'super': 10,
'unveiled': 1,
'notch': 1,
'hanging': 5,
'inventory': 4,
'article': 4,
'cooler': 4,
'current': 1,
. . . . .}
```

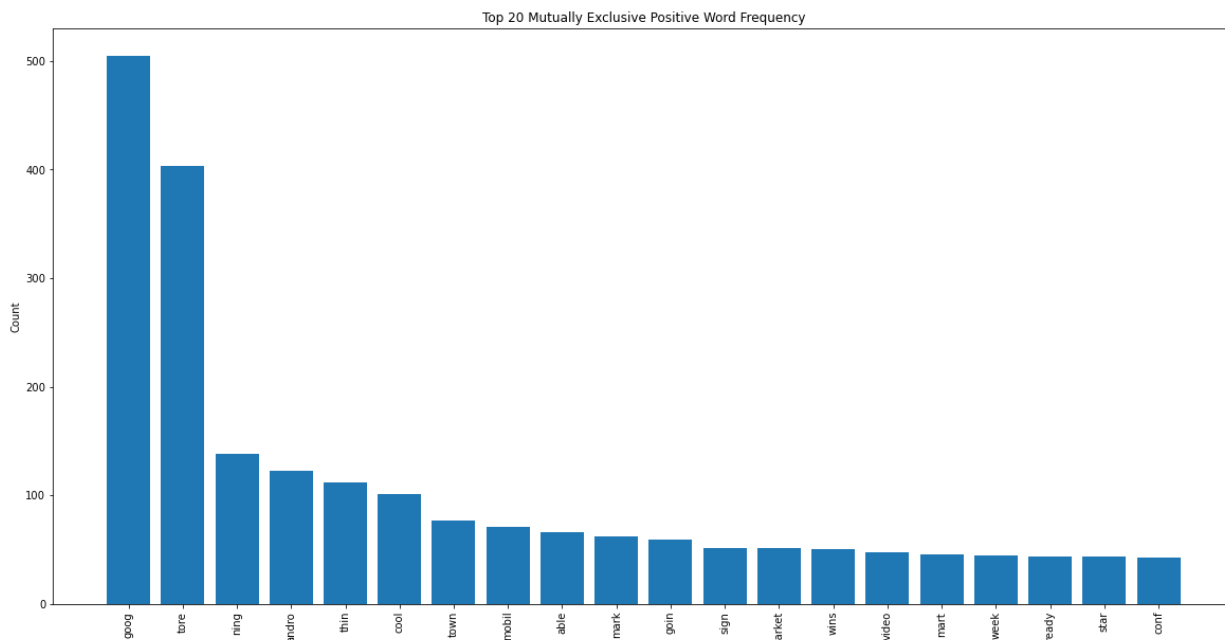
```
In [74]: 1 counts(diff_neg, negative['lemmed_text'])
```

started 08:23:58 2022-02-21, finished in 52ms

```
Out[74]: {'forgotten': 1,
'technical': 3,
'dense': 2,
'vuelta': 1,
'para': 3,
'gran': 2,
'diferencia': 1,
'revolution': 1,
'clumsily': 1,
'translated': 1,
'speaks': 1,
'watched': 3,
'staff': 1,
'facepalmed': 1,
'bandwidth': 1,
'tweeted': 1,
'dawdled': 1,
'blurs': 1,
'fades': 7,
'...' : ...}
```

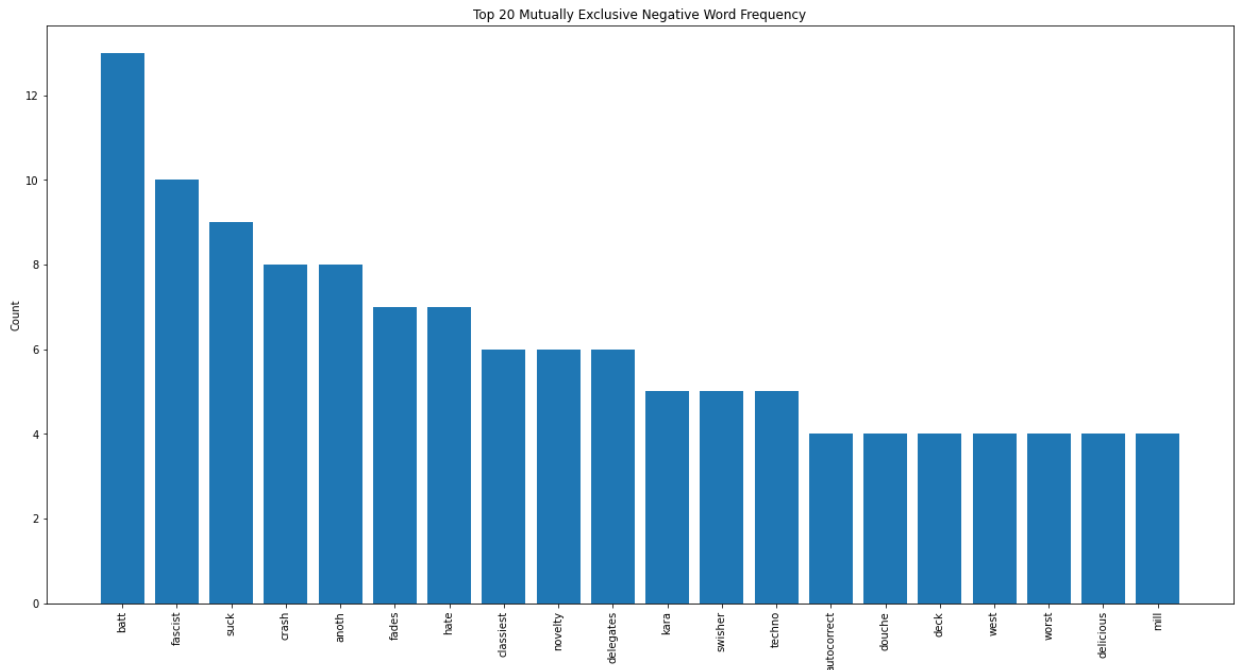
```
In [75]: 1 # View top 20 mutually exclusive positive words
2 positives_freq_dist = FreqDist(counts(diff_pos, positive['lemmed_text'])
3
4 visualize_top_20(positives_freq_dist, "Top 20 Mutually Exclusive Positi
```

started 08:23:58 2022-02-21, finished in 1.59s



```
In [76]: 1 # # View top 20 mutually exclusive negative words
2 negatives_freq_dist = FreqDist(counts(diff_neg, negative['lemmed_text'])
3
4 visualize_top_20(negatives_freq_dist, "Top 20 Mutually Exclusive Negati
```

started 08:24:00 2022-02-21, finished in 200ms



```
In [77]: 1 # Create a function to remove the mutual words and apply it to the train
2 def remove_samewords(token_list):
3     """
4     Given a list of tokens, return a list where the tokens
5     that are in both pos and neg have been removed
6     """
7     same_rmv_list = [token for token in token_list if token not in same
8     return same_rmv_list
9
10 pos_neg_train['sames_removed'] = pos_neg_train['lemmed_text_tokenized']
```

started 08:24:00 2022-02-21, finished in 1.15s

```
In [78]: 1 pos_neg_train.info()

started 08:24:01 2022-02-21, finished in 11ms

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2731 entries, 1435 to 4378
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   text                  2731 non-null   object
 1   text_tokenized        2731 non-null   object
 2   stopwords_removed     2731 non-null   object
 3   regex_text            2731 non-null   object
 4   regex_text_tokenized  2731 non-null   object
 5   stemmed_text          2731 non-null   object
 6   lemmmed_text          2731 non-null   object
 7   lemmmed_text_tokenized 2731 non-null   object
 8   target                2731 non-null   int64
 9   sames_removed         2731 non-null   object
dtypes: int64(1), object(9)
memory usage: 299.2+ KB
```

```
In [79]: 1 # Join the tokenized mutually exclusive word column for modeling
2 pos_neg_train['sames'] = [' '.join(text) for text in pos_neg_train.sames]

started 08:24:01 2022-02-21, finished in 5ms
```

```
In [80]: 1 # Split the dataset back into Training and testing data
2 # The only thing that should change is the X_train set
3 X_train_diff = pos_neg_train.drop(['target'], axis=1)
4 y_train_diff = y_train
5 X_test_diff = X_test
6 y_test_diff = y_test

started 08:24:01 2022-02-21, finished in 5ms
```

## Create a baseline model for the mutually exclusive data

```
In [81]: 1 # Vectorize our new training data
2 count_diff = CountVectorizer()
3 X_train_diff_vectorized = count_diff.fit_transform(X_train_diff['sames'])

started 08:24:01 2022-02-21, finished in 18ms
```

```
In [82]: 1 # Balance the data
2 smote_diff = SMOTE(k_neighbors=3)
3 X_train_diff_resampled, y_train_diff_resampled = smote_diff.fit_resample(
4     X_train_diff_vectorized, y_train_diff)

started 08:24:01 2022-02-21, finished in 10ms
```

```
In [83]: 1 # Build the model
2 baseline_model_diff = MultinomialNB()
3
4 baseline_model_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
5
6 # Evaluate the model
7 baseline_diff_cv = cross_val_score(baseline_model_diff, X_train_diff_re
8 print("Baseline_diff:", baseline_diff_cv.mean())
```


started 08:24:01 2022-02-21, finished in 19ms

Baseline\_diff: 0.670819007151536

```
In [84]: 1 # Test the model
2 X_test_diff_vectorized = count_diff.transform(X_test_diff['lemmed_text']
3 X_test_diff_preds = baseline_model_diff.predict(X_test_diff_vectorized)
4
5 print(classification_report(y_test_diff, X_test_diff_preds))
6
7 plot_confusion_matrix(baseline_model_diff, X_test_diff_vectorized, y_te
8 plt.grid(False)
9 plt.show()
```

started 08:24:01 2022-02-21, finished in 126ms

	precision	recall	f1-score	support
0	0.34	0.65	0.45	121
1	0.91	0.73	0.81	562
accuracy			0.71	683
macro avg	0.62	0.69	0.63	683
weighted avg	0.81	0.71	0.74	683



## And a bigram one too

```
In [85]: 1 count_diff_bigram = CountVectorizer(ngram_range=(1,2))
2 X_train_diff_bigram_vectorized = count_diff_bigram.fit_transform(X_train
```

started 08:24:01 2022-02-21, finished in 169ms


```
In [86]: 1 smote_diff_bigram = SMOTE(k_neighbors=3)
2 X_train_diff_resampled_bigram, y_train_diff_resampled_bigram = smote_diff_bigram.fit_resample(X_train_diff_bigram_vectorized, y_train_diff)
3
started 08:24:01 2022-02-21, finished in 8ms
```

```
In [87]: 1 baseline_model_diff_bigram = MultinomialNB()
2 baseline_model_diff_bigram.fit(X_train_diff_resampled_bigram, y_train_diff_resampled_bigram)
3
4 baseline_diff_bigram_cv = cross_val_score(baseline_model_diff_bigram, X_train_diff_resampled_bigram, y_train_diff_resampled_bigram, cv=5)
5 print("Baseline_diff : ", baseline_diff_cv.mean())
6 print("Baseline_diff_bigram : ", baseline_diff_bigram_cv.mean())
7
started 08:24:01 2022-02-21, finished in 18ms
```

```
Baseline_diff : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
```

```
In [88]: 1 X_test_diff_bigram_vectorized = count_diff_bigram.transform(X_test_diff_bigram_vectorized)
2 X_test_diff_bigram_preds = baseline_model_diff_bigram.predict(X_test_diff_bigram_vectorized)
3
4 print(classification_report(y_test_diff, X_test_diff_bigram_preds))
5
6 plot_confusion_matrix(baseline_model_diff_bigram, X_test_diff_bigram_vectorized, y_test_diff)
7 plt.grid(False)
8 plt.show()
9
started 08:24:01 2022-02-21, finished in 125ms
```

	precision	recall	f1-score	support
0	0.36	0.57	0.44	121
1	0.89	0.78	0.83	562
accuracy			0.74	683
macro avg	0.63	0.67	0.64	683
weighted avg	0.80	0.74	0.76	683



**Build a few more models with the mutually exclusive data and see how**

## they score

```
In [89]: 1 # Support Vector Machine
2 clf_diff = SVC()
3 clf_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
4 clf_diff_cv = cross_val_score(clf_diff, X_train_diff_resampled, y_train_diff_resampled, cv=5)
5 print("Baseline_diff      :", baseline_diff_cv.mean())
6 print("Baseline_diff_bigram :", baseline_diff_bigram_cv.mean())
7 print("SVM_diff           :", clf_diff_cv.mean())
```


started 08:24:02 2022-02-21, finished in 983ms

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
```

```
In [90]: 1 X_test_diff_preds_SVM = clf_diff.predict(X_test_diff_vectorized)
2
3 print(classification_report(y_test_diff, X_test_diff_preds_SVM))
4
5 plot_confusion_matrix(clf_diff, X_test_diff_vectorized, y_test_diff, cm=cm)
6 plt.grid(False)
7 plt.show()
```

started 08:24:03 2022-02-21, finished in 162ms

	precision	recall	f1-score	support
0	0.32	0.44	0.37	121
1	0.87	0.80	0.83	562
accuracy			0.74	683
macro avg	0.60	0.62	0.60	683
weighted avg	0.77	0.74	0.75	683



```
In [91]: 1 # K Nearest Neighbors (n=5)
2 knn_5_diff = KNeighborsClassifier(n_neighbors=5)
3 knn_5_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
4
5 knn_5_diff_cv = cross_val_score(knn_5_diff, X_train_diff_resampled, y_t
6 print("Baseline_diff      :", baseline_diff_cv.mean())
7 print("Baseline_diff_bigram:", baseline_diff_bigram_cv.mean())
8 print("SVM_diff           :", clf_diff_cv.mean())
9 print("KNN_5_diff         :", knn_5_diff_cv.mean())
```


started 08:24:03 2022-02-21, finished in 222ms

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
KNN_5_diff         : 0.612049033731842
```

```
In [92]: 1 knn_5_diff_preds = knn_5_diff.predict(X_test_diff_vectorized)
2
3 print(classification_report(y_test_diff, knn_5_diff_preds))
4
5 plot_confusion_matrix(knn_5_diff, X_test_diff_vectorized, y_test_diff,
6 plt.grid(False)
7 plt.show())
```

started 08:24:03 2022-02-21, finished in 185ms

	precision	recall	f1-score	support
0	0.21	0.92	0.35	121
1	0.94	0.27	0.42	562
accuracy			0.39	683
macro avg	0.58	0.60	0.39	683
weighted avg	0.81	0.39	0.41	683





```
In [93]: 1 # K Nearest Neighbors (n=3)
2 knn_3_diff = KNeighborsClassifier(n_neighbors=3)
3 knn_3_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
4
5 knn_3_diff_cv = cross_val_score(knn_3_diff, X_train_diff_resampled, y_t
6 print("Baseline_diff      :", baseline_diff_cv.mean())
7 print("Baseline_diff_bigram :", baseline_diff_bigram_cv.mean())
8 print("SVM_diff           :", clf_diff_cv.mean())
9 print("KNN_5_diff         :", knn_5_diff_cv.mean())
10 print("KNN_3_diff         :", knn_3_diff_cv.mean())
```


started 08:24:03 2022-02-21, finished in 236ms

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
KNN_5_diff         : 0.612049033731842
KNN_3_diff         : 0.5680500490008296
```

```
In [94]: 1 knn_3_diff_preds = knn_3_diff.predict(X_test_diff_vectorized)
2
3 print(classification_report(y_test_diff, knn_3_diff_preds))
4
5 plot_confusion_matrix(knn_3_diff, X_test_diff_vectorized, y_test_diff,
6 plt.grid(False)
7 plt.show())
```

started 08:24:03 2022-02-21, finished in 173ms

	precision	recall	f1-score	support
0	0.89	0.13	0.23	121
1	0.84	1.00	0.91	562
accuracy			0.84	683
macro avg	0.87	0.56	0.57	683
weighted avg	0.85	0.84	0.79	683



**Very odd, KNN\_3 positive accuracy was just about 100%, but only 13% on the negative class.**

This is the opposite of what it was doing on the KKN\_5 just above it, and on the lemmatized text

Predicted mostly negative

```
In [95]: 1 tree_diff = DecisionTreeClassifier()
2
3 tree_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
4
5 tree_diff_cv = cross_val_score(tree_diff, X_train_diff_resampled, y_train_diff_resampled, cv=5)
6 print("Baseline_diff      :", baseline_diff_cv.mean())
7 print("Baseline_diff_bigram :", baseline_diff_bigram_cv.mean())
8 print("SVM_diff           :", clf_diff_cv.mean())
9 print("KNN_5_diff          :", knn_5_diff_cv.mean())
10 print("KNN_3_diff          :", knn_3_diff_cv.mean())
11 print("Tree_diff           :", tree_diff_cv.mean())
```


started 08:24:04 2022-02-21, finished in 648ms

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
KNN_5_diff          : 0.612049033731842
KNN_3_diff          : 0.5680500490008296
Tree_diff           : 0.819464492612273
```

```
In [96]: 1 tree_diff_preds = tree_diff.predict(X_test_diff_vectorized)
2
3 print(classification_report(y_test_diff, tree_diff_preds))
4
5 plot_confusion_matrix(tree_diff, X_test_diff_vectorized, y_test_diff, c
6 plt.grid(False)
7 plt.show()
```

started 08:24:04 2022-02-21, finished in 119ms

	precision	recall	f1-score	support
0	0.29	0.66	0.41	121
1	0.90	0.66	0.76	562
accuracy			0.66	683
macro avg	0.60	0.66	0.58	683
weighted avg	0.79	0.66	0.70	683



```
In [97]: 1 # Random Forest (default)
2 forest_diff = RandomForestClassifier()
3
4 forest_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
5
6 forest_diff_cv = cross_val_score(forest_diff, X_train_diff_resampled, y
7 print("Baseline_diff      :", baseline_diff_cv.mean())
8 print("Baseline_diff_bigram :", baseline_diff_bigram_cv.mean())
9 print("SVM_diff           :", clf_diff_cv.mean())
10 print("KNN_5_diff         :", knn_5_diff_cv.mean())
11 print("KNN_3_diff         :", knn_3_diff_cv.mean())
12 print("Tree_diff          :", tree_diff_cv.mean())
13 print("Forest_diff        :", forest_diff_cv.mean())
14
```


started 08:24:04 2022-02-21, finished in 11.5s

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
KNN_5_diff         : 0.612049033731842
KNN_3_diff         : 0.5680500490008296
Tree_diff          : 0.819464492612273
Forest_diff        : 0.8166466511399448
```

```
In [98]: 1 forest_diff_preds = forest_diff.predict(X_test_diff_vectorized)
2
3 print(classification_report(y_test_diff, forest_diff_preds))
4
5 plot_confusion_matrix(forest_diff, X_test_diff_vectorized, y_test_diff,
6 plt.grid(False)
7 plt.show())
```

started 08:24:16 2022-02-21, finished in 289ms

	precision	recall	f1-score	support
0	0.30	0.67	0.42	121
1	0.90	0.67	0.77	562
accuracy			0.67	683
macro avg	0.60	0.67	0.59	683
weighted avg	0.80	0.67	0.71	683



```
In [99]: 1 # Random Forest (bootstrap=False)
2 forest_boot_diff = RandomForestClassifier(bootstrap=False)
3
4 forest_boot_diff.fit(X_train_diff_resampled, y_train_diff_resampled)
5
6 forest_boot_diff_cv = cross_val_score(forest_boot_diff, X_train_diff_re
7 print("Baseline_diff      :", baseline_diff_cv.mean())
8 print("Baseline_diff_bigram :", baseline_diff_bigram_cv.mean())
9 print("SVM_diff           :", clf_diff_cv.mean())
10 print("KNN_5_diff          :", knn_5_diff_cv.mean())
11 print("KNN_3_diff          :", knn_3_diff_cv.mean())
12 print("Tree_diff           :", tree_diff_cv.mean())
13 print("Forest_diff         :", forest_diff_cv.mean())
14 print("Forest_diff_boot    :", forest_boot_diff_cv.mean())
15
```


started 08:24:16 2022-02-21, finished in 19.6s

```
Baseline_diff      : 0.670819007151536
Baseline_diff_bigram : 0.6162036460377482
SVM_diff           : 0.8860094993454805
KNN_5_diff          : 0.612049033731842
KNN_3_diff          : 0.5680500490008296
Tree_diff           : 0.819464492612273
Forest_diff         : 0.8166466511399448
Forest_diff_boot    : 0.8266160285591406
```

```
In [100]: 1 forest_boot_diff_preds = forest_boot_diff.predict(X_test_diff_vectorize
2
3 print(classification_report(y_test_diff, forest_boot_diff_preds))
4
5 plot_confusion_matrix(forest_boot_diff, X_test_diff_vectorized, y_test_
6 plt.grid(False)
7 plt.show()
```

started 08:24:36 2022-02-21, finished in 304ms

	precision	recall	f1-score	support
0	0.31	0.64	0.42	121
1	0.90	0.70	0.79	562
accuracy			0.69	683
macro avg	0.61	0.67	0.60	683
weighted avg	0.80	0.69	0.72	683



**Let's Check if TF-IDF can do any better on mutually exclusive words**

```
In [101]: 1 tfidf_diff = TfidfVectorizer(max_features=60)
2
3 X_train_diff_tfidf_vectorized = tfidf_diff.fit_transform(X_train_diff['
```

started 08:24:36 2022-02-21, finished in 18ms

```
In [102]: 1 smote_diff_tfidf = SMOTE(k_neighbors=3)
2 X_train_diff_resampled_tfidf, y_train_diff_resampled_tfidf = smote_diff
3 X_train_diff_tfidf_vectorized, y_train_diff)
```

started 08:24:36 2022-02-21, finished in 10ms

```
In [103]: 1 baseline_model_diff_tdidf = MultinomialNB()
2
3 baseline_model_diff_tdidf.fit(X_train_diff_resampled_tdidf, y_train_diff)
4
5 # Evaluate the model
6 baseline_diff_tdidf_cv = cross_val_score(baseline_model_diff_tdidf, X_test_diff, y_test_diff, cv=5)
7 print("Baseline_diff_tdidf:", baseline_diff_tdidf_cv.mean())
```

started 08:24:36 2022-02-21, finished in 20ms

Baseline\_diff\_tdidf: 0.5049990246837273

```
In [104]: 1 X_test_diff_tdidf_vectorized = tfidf_diff.transform(X_test_diff['lemmed'])
2 X_test_diff_tdidf_preds = baseline_model_diff_tdidf.predict(X_test_diff_tdidf_vectorized)
3
4 print(classification_report(y_test_diff, X_test_diff_tdidf_preds))
```

started 08:24:36 2022-02-21, finished in 13ms

	precision	recall	f1-score	support
0	0.19	0.91	0.31	121
1	0.89	0.16	0.27	562
accuracy			0.29	683
macro avg	0.54	0.54	0.29	683
weighted avg	0.77	0.29	0.28	683

No, it cannot

## Results

While using mutually exclusive words in a Random Forest model showed some accuracy on predictions (negative: 65%, positive: 70%), It was not more accurate than our baseline model using Naive Bayes on the full text with lemmatization. Count Vectorization also proved to be more accurate than using TF-IDF.

## Final Model:

Multinomial Naive Bayes:  
Lemmatized text with Count Vectorization

Positive prediction accuracy: 88%  
Negative prediction accuracy: 64%

## Interpretation:



Imbalanced data posed problems when testing the model

- Predicted the true positive class much better than the true negative class leading to a higher false positive rate

Words expressing positive sentiment like "great", and "awesome" appeared in both classes

- Could lead to false negatives
- Did not appear in the mutually exclusive word models

## Conclusions

- **Based on a general sense of the dataset the overwhelming sentiment was positive.**
- **This positive sentiment can be used in an overall investment strategy of emerging technologies**
- **This was a small dataset and more data is needed to improve the model**

- The use of synthetically produced data was used to balance the dataset
- With more data a downsampling technique could be employed instead

## Further Research

- Gather more data to provide more balance for testing
- Explore Neural Networking techniques and Word2Vec for modeling
- Explore techniques for weighting specific words
- Look at sentiment on an individual product level

In [ ]:

1